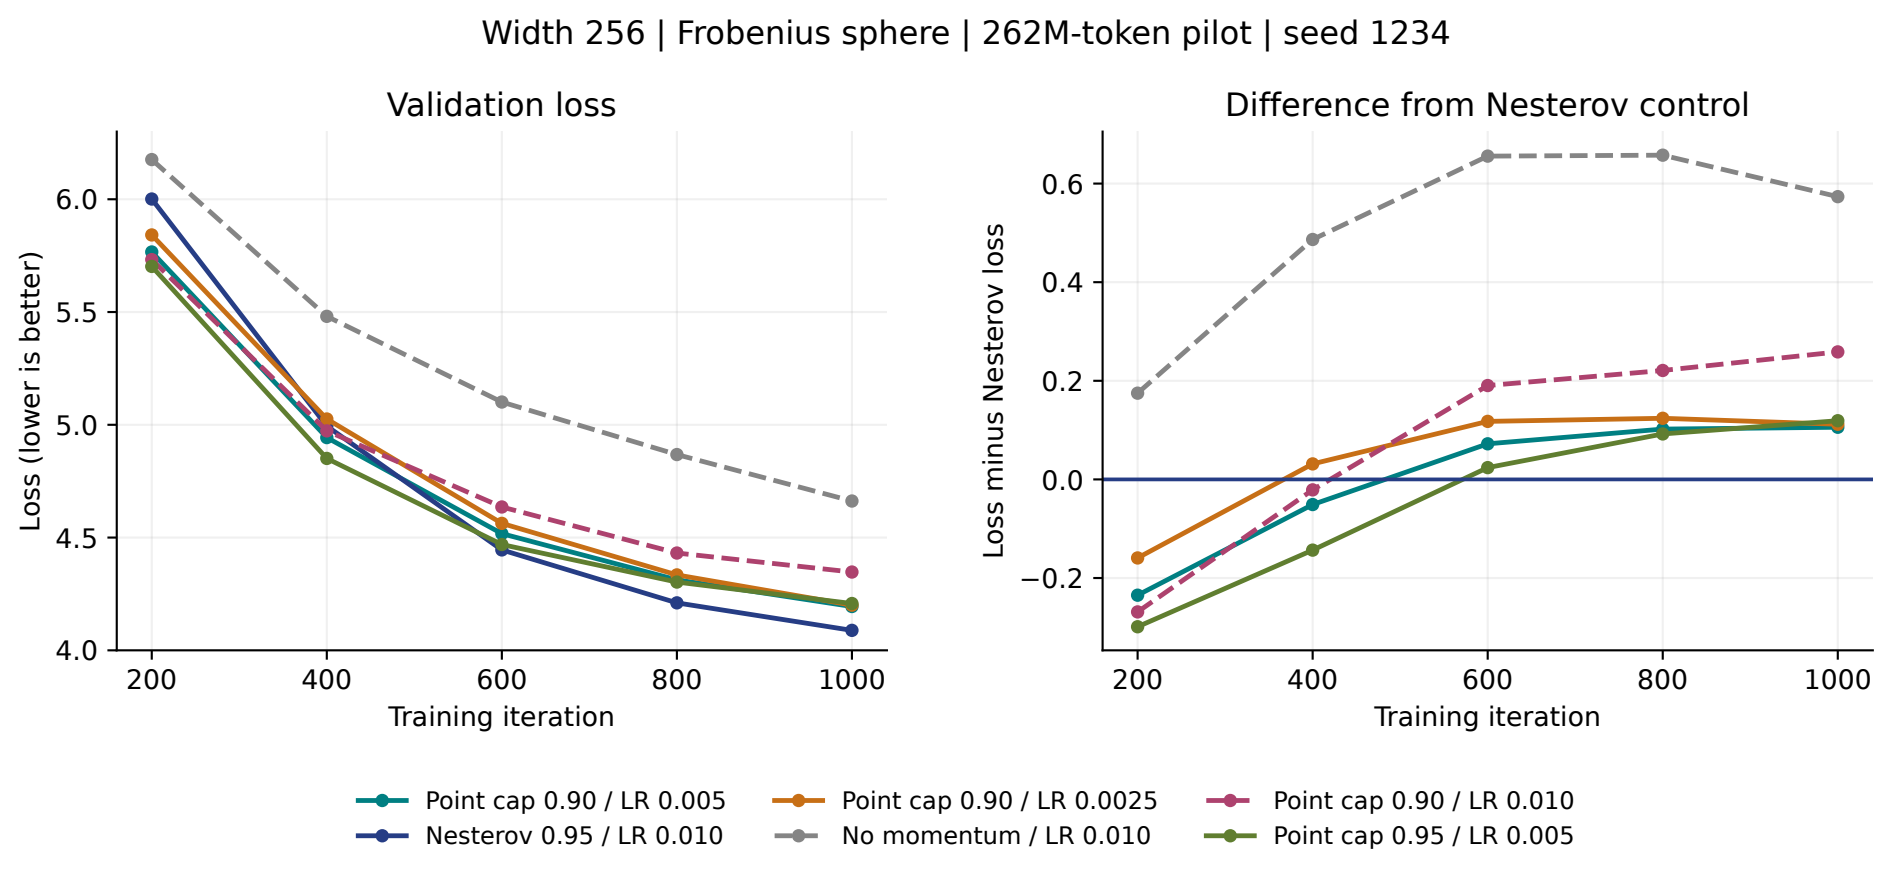
Width 256 | Frobenius sphere | 262M-token pilot | seed 1234
Validation loss
Loss (lower is better)
6.0
5.5
5.0
4.5
4.0
200
400
600
800
1000
Training iteration
Difference from Nesterov control
Loss minus Nesterov loss
0.6
0.4
0.2
0.0
−0.2
200
400
600
800
1000
Training iteration
Point cap 0.90 / LR 0.005
Point cap 0.90 / LR 0.0025
Point cap 0.90 / LR 0.010
Nesterov 0.95 / LR 0.010
No momentum / LR 0.010
Point cap 0.95 / LR 0.005